

Engineering Toward N: Cambridge Semantics as Empirical Validation of the Tholonic N-D-C Framework

J. W. Milton

Clarity Coalition

16 April 2026

v1.0

Keywords: Cambridge Semantics; N-D-C framework; empirical validation; knowledge graphs; consilience; engineering toward N

Abstract

The tholonic framework models any coherent system as a recursive triadic structure in which a stable emergent state N (negotiation) arises from the balance of a bounding parameter D (definition) and an integrating parameter C (contribution). Prior papers in this series establish the mathematical grounding of this framework in prime recursion (paper 1), apply it to supply chain transparency scoring (paper 2), prove the structural necessity of exactly three roles (paper 3), and examine its game-theoretic and geometric extensions (papers 4 and 5). These applications are theoretical: the framework has been instantiated by design.

This paper presents a different class of evidence: an independent empirical convergence. Cambridge Semantics, a knowledge graph company founded in 2007 and acquired by Altair (2024) and subsequently Siemens (2025), spent two decades engineering large-scale semantic data fabrics. Their architecture, built entirely from engineering necessity and validated against production deployments handling hundreds of billions of RDF triples, exhibits a precise structural correspondence to the tholonic N-D-C hierarchy across every level: ontology as D , data as C , knowledge graph as N ; the $D \approx C$ balance principle as the discovered condition for ontology stability; opacity as a structural break in the parent-to-child- N chain; entity explosion as C outrunning D ; and provenance as a property of the recursive hierarchy rather than a metadata add-on.

Cambridge Semantics discovered these structural facts without any knowledge of the tholonic framework. The convergence is therefore not an analogy constructed after the fact. It is an independent rediscovery of the same underlying structure by practitioners engineering toward a goal they could not formally name.

We further show that the relationship is bidirectional: Cambridge Semantics' empirical methodology provides the tholonic model with a construction sequence (bottom-up ontology lifecycle), a philosophy of provisional N states, and a proof-of-concept implementation architecture (Cognee, as of May 2026). The paper concludes with a summary

table of open gaps in both systems and a statement of the most important remaining formal problem: recursive SPARQL-based tholonic chain traversal at the entity level.

1 Introduction

1.1 The problem with purely theoretical frameworks

The tholonic N-D-C framework has been argued from mathematical first principles across five prior papers. Papers 1 through 6 of this series establish that the framework is internally consistent, that its constants emerge naturally from prime recursion, and that it maps coherently onto supply chain scoring, game theory, and twistor geometry. What these papers do not establish is whether the framework describes something real about how complex systems actually behave, or whether it is an elegant formal construction whose applicability to the world is limited.

The standard response to this concern is empirical: find a domain where the framework's predictions can be tested against observed behavior. This paper takes a different approach. Rather than testing predictions, we identify an independent engineering effort that, over two decades and under the pressure of production failures and successes at enterprise scale, converged on the same structural architecture that the tholonic framework predicts from first principles.

The convergence is not a proof. It is a form of consilience: multiple independent lines of evidence pointing to the same structure, where the independence of the sources rules out the possibility that the correspondence is an artifact of shared assumptions.

1.2 Cambridge Semantics as evidence

Cambridge Semantics was a knowledge graph company that built production systems for some of the most data-complex domains in enterprise computing: clinical trial integration across fifteen siloed research centers, insider trading surveillance integrating trading records with communications and physical access logs, pharmaceutical adverse event pipelines processing hundreds of thousands of reports per year. Their architecture, AnzoGraph and Anzo, handled tens to hundreds of billions of RDF triples in production with clusters benchmarked against a trillion triples.

They discovered, through failure and iteration, that robust systems require three things to be in balance: a formal constraint structure that gives meaning (the ontology), a flow of real-world facts (the data), and a stable emergent synthesis of the two (the knowledge graph). They discovered that over-specifying the constraint structure before the data arrives causes the synthesis to fail. They discovered that loading data without adequate constraint structure produces an incoherent and unusable result. They discovered that the stability of the synthesis is always provisional, always dependent on the current state of both constraint and data.

These are not independently motivated philosophical conclusions. They are engineering findings, extracted from the practical experience of building systems that could not be made to work any other way.

1.3 Structure of this paper

Section 2 recaps the relevant structure of the tholonic framework. Section 3 describes the Cambridge Semantics architecture from primary sources. Section 4 develops the structural alignment in detail. Section 5 reverses the direction and asks what Cambridge Semantics offers the tholonic model. Section 6 distinguishes the domains of the two systems and shows why they are complementary rather than competing. Section 7 reports confirmed implementation status as of May 2026. Section 8 states the open problems and summarizes the contribution.

2 The Tholonic Framework: Relevant Structure

The Tholonic N-D-C framework models any coherent system as a recursive triadic structure:

- **N (Negotiation):** the stable, coherent instantiation at a given level. It is not directly measured. It emerges from the balance of D and C . It is simultaneously the product of the level above and the source for the level below.
- **D (Definition):** constraints, limitations, boundaries, and specifications. Defines what a thing IS. Internally focused.
- **C (Contribution):** outputs, applications, connections, and integrations. Defines what a thing DOES. Externally focused.

The full recursive cycle:

Parent $N \rightarrow$ differentiates into D and $C \rightarrow D$ and C negotiate $\rightarrow$ Child N instantiates

Child $N \rightarrow$ becomes next Parent N

The sustainability principle: systems are most stable and efficient when $D \approx C$. Imbalance in either direction increases energy cost and degrades the N state. When $D \gg C$, the system is over-constrained: the negotiation space is so narrow that no real-world data can instantiate a valid N . When $C \gg D$, the system is under-constrained: data flows without sufficient definition to produce a coherent synthesis.

The mathematical grounding: when the first three prime numbers (2, 3, 5) are assigned to N, D, and C respectively and the recursive model is applied, the fundamental constants emerge naturally: π , ϕ (golden ratio), $\sqrt{2}$, e , and $\ln 2$. The framework is grounded in the irreducible relationships between the first primes.

Paper 6 of this series shows that the assignment of roles to positions is not arbitrary: the cardinality-1, cardinality-2, and cardinality-3 structures each exhibit a qualitative transition at the corresponding integer that is isomorphic to the $N \rightarrow D \rightarrow C$ role transition, across five independent mathematical domains and independently recovered by Kant (Unity, Plurality, Totality), Peirce (Firstness, Secondness, Thirdness), and Spencer-Brown (distinction, marked state, re-entry). The role assignment follows from what the integers structurally are.

3 The Cambridge Semantics Architecture

3.1 The founding problem

Cambridge Semantics was founded in 2007 by engineers who had spent years at IBM wrestling with a problem that defeated relational databases: how do you integrate radically heterogeneous enterprise data when the data never stops changing and you cannot predict all the questions users will ask?

The relational database model assumes a fixed schema defined upfront. It fails when entity types proliferate beyond what any fixed schema can contain (what the founders called “entity explosion”), when entities have complex subclass relationships, when data evolves continuously, and when users need to ask questions nobody anticipated. Their insight was that the relational model forces storage optimized for the machine rather than for actual knowledge structure.

3.2 The technical solution

Their answer was the semantic knowledge graph based on W3C open standards.

RDF (Resource Description Framework) represents every fact as a triple: subject-predicate-object. Every entity is a node. Every relationship is an edge. The enterprise data universe becomes a connected graph of triples.

OWL (Web Ontology Language) provides the ontology: an abstract schema that describes what everything in the graph means. The ontology travels with the data, making the data self-describing independent of any originating application.

SPARQL is the graph traversal query language: it lets you ask questions of the entire integrated fabric without knowing upfront where the answer lives.

Sean Martin, Cambridge Semantics CTO, summarized the architecture in one sentence: “The knowledge graph is the ontology with the data.”

3.3 The scale breakthrough

Their critical engineering achievement was AnzoGraph: a massively parallel, in-memory graph OLAP database. Production deployments handled tens to hundreds of billions of RDF triples. Major use cases included insider trading surveillance (30 billion triples per year with 3-year rolling windows), FDA drug data integration, and pharmaceutical adverse event pipelines (quarter-million reports per year). Altair Engineering acquired Cambridge Semantics in April 2024; Siemens acquired Altair in March 2025.

4 The Structural Alignment

4.1 The direct mapping

Sean Martin’s one-sentence definition of the knowledge graph is a direct tholonic statement. “The knowledge graph is the ontology with the data” maps precisely to “ N is D with C .”

Cambridge Semantics	Tholonic role	Reason
Ontology (OWL)	<i>D</i> (Definition)	Defines classes, properties, constraints, hierarchies: what everything IS. Does not flow. Limits, specifies, constrains.
Data (RDF triples)	<i>C</i> (Contribution)	Real-world facts, relationships, measurements, events: what the enterprise DOES and PRODUCES. Flows, connects, integrates.
Knowledge graph	<i>N</i> (Negotiation)	The emergent, stable, coherent synthesis. Remove the ontology and the data is meaningless. Remove the data and the ontology is an empty schema.

Cambridge Semantics discovered this empirically across two decades of engineering. The tholonic model names it formally.

4.2 The recursive hierarchy

The tholonic model specifies that each child *N* becomes the parent *N* for the next level down. Cambridge Semantics built exactly this three-level recursive structure:

Level 1: Enterprise data fabric. Parent *N* is the enterprise knowledge graph. *D* is the upper ontology (canonical business concepts, their definitions and relationships). *C* is all data flows from all source systems. Child *N* is the Graphmart: the negotiated, activated instance of a specific analytic domain.

Level 2: The Graphmart. Parent *N* is the Graphmart itself. *D* is the data layer definitions, transformation rules, SPARQL integration queries, and access control policies. *C* is actual data loaded from specific sources through specific pipelines. Child *N* is the activated, queryable in-memory graph.

Level 3: The data layer. Parent *N* is the data layer definition. *D* is source schema mappings, ontology mappings, step sequences, and validation rules. *C* is raw data ingested from the specific source system. Child *N* is the sub-graph contributed to the graphmart.

Three levels of the tholonic hierarchy, all present and operating in Cambridge Semantics' architecture, with child *N* becoming parent *N* at each step down. They built this structure because it works. The tholonic model explains why it works.

4.2.1 4.3 The $D \approx C$ sustainability discovery

Cambridge Semantics' most empirically validated finding was this: ontologies designed before the data ($D \gg C$) almost always fail. Martin stated it directly: "We find it's very

important to maintain flexibility and essentially model the data you actually have, then use upper ontology as a kind of highway map to the concepts. You get these very low-level representations bubbling up out of the data and then you try to abstract them as much as you can.”

This is the tholonic $D \approx C$ balance principle, stated in engineering terms.

- When $D \gg C$: over-specified ontology, rigid upfront schema. The system becomes brittle. The N that “emerges” does not correspond to any real stable state in the data. It is a formal constraint imposed on reality rather than emerging from it.
- When $C \gg D$: raw data ingestion with no ontological structure. The result is an unusable pile of triples with no coherent meaning. N has no definition to stabilize it.
- When $D \approx C$: the ontology is expressive enough to give meaning, flexible enough to accommodate the actual shape of real data. N is stable and productive.

Cambridge Semantics’ customers who succeeded followed this pattern. Their customers who failed tended to either over-specify ontologies upfront or load massive data with inadequate semantic modeling.

4.2.2 4.4 Entity explosion as C outrunning D

Cambridge Semantics identified “entity explosion” as the core failure mode of relational models: too many entity types and too many complex relationships for a fixed schema to contain. In tholonic terms this is precise: entity explosion is C outrunning D . The contribution space (the actual complexity of real-world entities and relationships) grows faster than the definition structure (the relational schema) can classify and constrain it. The negotiation collapses because there is no stable N that can emerge from an overloaded C with an insufficient D .

Their solution, semantic graphs with flexible OWL ontologies, is exactly a D designed to expand to meet the C . OWL ontologies can be extended with new classes and properties without breaking existing queries. This is D engineered for $D \approx C$ balance, which is the structural reason it works where relational models fail.

4.3 Opacity as structural chain break

Cambridge Semantics encountered many “opaque” domains: data sources they could not integrate because source system owners refused access, or because transformation logic was proprietary. They treated opacity as a practical obstacle.

The tholonic framework elevates opacity to a structural finding. A phase where the parent- N to child- N transition cannot be traced is not merely an obstacle to work around. It is a structural break in the tholonic hierarchy, diagnostically significant in itself.

Opacity occurs when C cannot be constrained by D (the data source refuses to be described), when D cannot be populated by C (the source refuses to provide data), or when the parent N governing the negotiation is itself inaccessible (the owning authority refuses to participate at all). These are structurally distinct situations that require structurally distinct responses, something Cambridge Semantics did not formally distinguish.

4.4 Provenance as hierarchical chain property

Martin identified provenance as a pressing unsolved problem: “Keeping track of what was established by a human versus what was inferred, and which models and training data were used.” As knowledge graphs grow through automated inference and LLM-based extraction, facts appear whose lineage is opaque.

The tholonic framework solves this structurally. Because each N can be traced to its parent D and C , and each of those to the N above them, provenance is not a metadata problem to be bolted on after the fact. It is a structural property of the hierarchy. Every triple has a position in the recursive chain: which ontological constraint (D) permitted it to exist, which data flow (C) produced it, and which parent N authorized the negotiation space in which it emerged.

5 What Cambridge Semantics Offers the Tholonic Model

The previous section asked how the tholonic framework could enhance Cambridge Semantics. This section inverts the question.

5.1 The bottom-up construction sequence

Cambridge Semantics discovered empirically that the correct sequence for building ontologies is not to define everything upfront and then match data to it. Start by ingesting real data, let the actual shape of that data define low-level concepts, and then progressively abstract those concepts upward toward a stable upper ontology.

Translated into tholonic terms: when building a new tholonic hierarchy for an unfamiliar domain, do not start by defining the parent N . Start by observing C (what is actually produced and flowing), let the D constraints surface from the boundaries that actually govern those flows, and allow the N to emerge from their negotiation. Only then attempt to connect that bottom-level N to a higher-level N .

The tholonic framework describes the structure of a hierarchy. It does not specify how to build one. Cambridge Semantics provides that construction sequence.

5.1.1 5.2 The provisional N concept

AnzoGraph can load unstructured or schema-less data directly into a graph without pre-mapping, and then apply transformation layers on top to shape it into meaningful structures. This suggests that the tholonic model benefits from a concept of a “provisional N ”: an entity that appears to be a stable emergent point but has not yet had its D and C fully formalized. A provisional N can be treated as a working hypothesis, refined as more D constraints are identified and more C flows are observed. This allows tholonic modeling to begin before the structure of a domain is fully understood.

5.2 The metadata-as-data principle

Cambridge Semantics made all metadata (schemas, mappings, transformation rules, access policies, provenance records) into first-class citizens of the same graph as the data itself. The implication for the tholonic model: the description of the hierarchy should be IN the hierarchy itself. A hierarchy that requires an external document to explain its own structure violates the same principle that makes Cambridge Semantics' approach robust.

5.3 The living model philosophy

Cambridge Semantics concluded that a knowledge graph is never complete. It is a continuously evolving artifact. New C flows in. New D constraints are discovered. The N at any given level is always provisional, always subject to refinement.

The tholonic model implicitly treats N as a stable resolved state. Cambridge Semantics challenges this framing. N is stable relative to the D and C present at a given moment. As D and C evolve, N must renegotiate. The “stability” of N is not permanence. It is coherence under current conditions. Every tholonic hierarchy is always in process, always a snapshot of an ongoing negotiation.

6 Domain Distinction: Why the Two Systems Are Complementary

Despite the structural overlap, Cambridge Semantics and the tholonic model as applied in the cTVF project target fundamentally different domains.

Cambridge Semantics was built to solve problems inside large organizations. The unit of analysis was the individual instance: a specific patient, a specific transaction, a specific regulatory submission. The knowledge graph was pointed inward: what does our organization know about this specific case?

The cTVF application of the tholonic model operates at the phase level of entire supply chains across multiple custodians and jurisdictions. The unit of analysis is the phase: a discrete physical transformation or custody change in a multi-step supply chain. The relevant measures are physical flow rates, custody transitions, efficiency gaps across phases, opacity classification, and predictive modeling of end-to-end behavior. The question is: what is the structural integrity of this supply chain as a whole system, where are the stress points, and what does the phase-level data predict about downstream behavior?

In Cambridge Semantics' deployments, a node is typically an instance. The graph is dense with individuals. In the cTVF model, a node is typically a phase-level aggregate. The graph is sparse in individuals but dense in structural relationships between abstract supply chain concepts.

These are not competing applications. They are complementary. The cTVF tholonic model provides the phase-level structural map: where are the bottlenecks, which phases are

opaque, what does the D - C balance predict about system stability. A Cambridge Semantics-style knowledge graph would provide the instance-level detail within any single phase: the specific custodians, the specific transactions, the specific compliance records.

In tholonic terms: cTVF operates at the level of the Phase Map (child N at Level 1). Cambridge Semantics operates inside each individual phase (child N at Level 2). Both are necessary. Neither replaces the other.

7 Implementation Evidence (May 2026)

As of May 2026, the tholonic framework has a working implementation vehicle: Cognee, an open-source knowledge engine combining vector search, graph storage, and LLM context grounding. A local deployment extends Cognee with a Tholonic semantic layer.

The Tholonic knowledge graph ontology is live under the IRI <http://tholonia.org/ontology/merged-> declaring 590+ `owl:Class` entries under the `thol:` prefix, merging three knowledge graph collections: TVF modeling, the Tholonia theoretical corpus, and I Ching trigram structure. Every Tholonic concept is a first-class RDF citizen with a stable IRI, a label, and graph traversal access.

The Tholonic N-D-C framework has been formally expressed in OWL. A `tholonic_ndc` abstract property group defines child `owl:ObjectProperty` entries including `has_ndc_role`, `has_n_state`, `has_d_parameter`, `has_c_parameter`, `has_balance_score`, `has_coupling_ratio`, `exhibits_self_similarity`, and `recursive_architecture`. These predicates carry formal semantics and subproperty hierarchies and can be reasoned over by any OWL-compliant tool. The `has_balance_score` and `has_coupling_ratio` properties create a direct hook for the quantitative $D \approx C$ balance metrics: the hook exists; populating it with computed values is the remaining step.

SPARQL is implemented at the ontology layer, querying the formal TTL files via `rdflib` and caching results in PostgreSQL. What is not yet implemented is SPARQL for recursive entity-level chain traversal using `+/*` property path operators. The ontology is well-formed and ready; a SPARQL endpoint over it would close this gap without changing the cache architecture.

8 Summary and Open Problems

Cambridge Semantics built a machine for instantiating N states at enterprise scale, without a theory of what N is or why the machine works. The tholonic model provides a theory of what N is and why the machine works, without a production-proven implementation path. Together they constitute something closer to a complete system.

Dimension	Cambridge Semantics Provides	Tholonic Framework Provides	Status (May 2026)
Formal structure	Absent (empirical only)	Present (N-D-C recursive hierarchy)	Implemented in OWL
Mathematical grounding	Absent	Present (ϕ , π , e from prime recursion)	Not yet wired to computed metrics
Production implementation	Present (Anzo, AnzoGraph)	Absent	Implemented via Cognee
RDF encoding of tholonics	Present (methodology)	Absent	Done: 590+ classes, http://tholonia.org/ontology/
OWL ontology of N-D-C	Present (methodology)	Absent	Done: <code>tholonic_ndc</code> property group
SPARQL query layer	Present	Absent	Partial: ontology layer done; entity recursive traversal not yet
Balance metrics ($D \approx C$)	Absent	Present (principle)	Properties exist; computation not wired
Opacity classification	Informal	Formal (structural break in chain)	Not yet implemented
Provenance	Partially solved	Fully structural (chain traceability)	Partial: lineage fields present
Construction sequence	Present (bottom-up lifecycle)	Absent	Adopted informally in KG build process
Self-description	Present (metadata-as-data)	Absent	Partial: metadata in relational cache, not in graph
Living model philosophy	Present	Absent (N treated as stable)	Implemented: Cognee continuously ingests
Propagation rules	Absent	Present	Not yet implemented

The most important open formal problem identified by this analysis is recursive tholonic chain traversal at the entity level: given any entity, return all ancestors in the tholonic hierarchy up to the root N ; given a D constraint, find all entities it governs; identify all entities where the D - C balance score falls outside a specified range. These cross-cutting structural queries require the $+/^*$ path operators in SPARQL operating against a combined graph of OWL ontology and entity data. Closing this gap would make the tholonic hierarchy fully queryable as a formal structure rather than only as a cached flat triple store.

A second open problem is the computation and validation of the $D \approx C$ balance metric for a concrete graph domain. The formal hook exists. The philosophical argument for why ϕ governs the optimal balance point in self-similar recursive systems is made in paper 1. What is missing is an empirical study: take a production knowledge graph, compute the

D - C balance at each level, and verify that the stability and query performance of the graph correlates with proximity to the predicted balance point.

9 References

- **Milton 2026a**: Milton, J. W. “Recursive Tholonic Five Constants.” 2026. (Paper 1 of this series.)
- **Milton 2026b**: Milton, J. W. “Supply Chain Transparency via Phase-resolved Classification and Indexing (TVPCI).” 2026. (Paper 2 of this series.)
- **Milton 2026c**: Milton, J. W. “Minimal Recursive Triadic Framework.” 2026. (Paper 3 of this series.)
- **Milton 2026d**: Milton, J. W. “Game-Theoretic Triadic Balance.” 2026. (Paper 4 of this series.)
- **Milton 2026e**: Milton, J. W. “Tholonic Twistor Connection.” 2026. (Paper 5 of this series.)
- **Milton 2026f**: Milton, J. W. “The Qualitative Nature of One, Two, and Three: Structural Role Assignment in Minimal Recursive Systems.” 2026. ([docnav/Research/papers/6_qualita](#))
- **Kant 1781**: Kant, I. *Critique of Pure Reason*. Translated by N. K. Smith. Macmillan, 1929 [1781]. (A80/B106: Table of categories; categories of quantity: Unity, Plurality, Totality.)
- **Peirce 1931**: Peirce, C. S. *Collected Papers of Charles Sanders Peirce*. Vol. 1. Harvard University Press, 1931. (§§300-353: Firstness, Secondness, Thirdness; §346: formal irreducibility of Thirdness.)
- **Spencer-Brown 1969**: Spencer-Brown, G. *Laws of Form*. Allen and Unwin, 1969.
- **Martin 2022**: Martin, S. “Cambridge Semantics CTO Interview.” Earley AI Podcast, 2022.
- **Szekely 2020**: Szekely, B. “Can Graph Integrate Data At Scale?” Medium, 2020.